

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import KFold, StratifiedKFold
from sklearn.metrics import roc_auc_score, accuracy_score, confusion_matrix
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.simplefilter(action='ignore', category=UserWarning)
```

```
In [2]: df=pd.read_csv("churn_prediction.csv")
```

```
In [3]: df
```

```
Out[3]:
```

	customer_id	vintage	age	gender	dependents	occupation	city	customer_nw
0	1	3135	66	Male	0.0	self_employed	187.0	
1	2	310	35	Male	0.0	self_employed	NaN	
2	4	2356	31	Male	0.0	salaried	146.0	
3	5	478	90	NaN	NaN	self_employed	1020.0	
4	6	2531	42	Male	2.0	self_employed	1494.0	
...	...	...	...	...	...	...	...	...
28377	30297	1845	10	Female	0.0	student	1020.0	
28378	30298	4919	34	Female	0.0	self_employed	1046.0	
28379	30299	297	47	Male	0.0	salaried	1096.0	
28380	30300	2585	50	Male	3.0	self_employed	1219.0	
28381	30301	2349	18	Male	0.0	student	1232.0	

28382 rows × 21 columns

```
In [4]: df.isnull().sum()
```

```
Out[4]: customer_id      0
vintage      0
age      0
gender      525
dependents    2463
occupation    80
city      803
customer_nw_category      0
branch_code      0
days_since_last_transaction    3223
current_balance      0
previous_month_end_balance      0
average_monthly_balance_prevQ      0
average_monthly_balance_prevQ2      0
current_month_credit      0
previous_month_credit      0
current_month_debit      0
previous_month_debit      0
current_month_balance      0
previous_month_balance      0
churn      0
dtype: int64
```

```
In [5]: df['gender'].value_counts()
```

```
Out[5]: gender
Male      16548
Female    11309
Name: count, dtype: int64
```

```
In [6]: #here converting male to 1 and female as 0 and filling nul values as -1

dict_gender={'Male':1,'Female':0}
df.replace({'gender':dict_gender}, inplace=True)
df['gender']= df['gender'].fillna(-1)
```

In [7]: `df`

Out[7]:

	customer_id	vintage	age	gender	dependents	occupation	city	customer_nw
0	1	3135	66	1.0	0.0	self_employed	187.0	
1	2	310	35	1.0	0.0	self_employed	NaN	
2	4	2356	31	1.0	0.0	salaries	146.0	
3	5	478	90	-1.0	NaN	self_employed	1020.0	
4	6	2531	42	1.0	2.0	self_employed	1494.0	
...	...	...	...	...	...	...	...	...
28377	30297	1845	10	0.0	0.0	student	1020.0	
28378	30298	4919	34	0.0	0.0	self_employed	1046.0	
28379	30299	297	47	1.0	0.0	salaries	1096.0	
28380	30300	2585	50	1.0	3.0	self_employed	1219.0	
28381	30301	2349	18	1.0	0.0	student	1232.0	

28382 rows × 21 columns

In [8]: `df['dependents'].value_counts()`

Out[8]:

dependents	count
0.0	21435
2.0	2150
1.0	1395
3.0	701
4.0	179
5.0	41
6.0	8
7.0	3
9.0	1
52.0	1
36.0	1
50.0	1
8.0	1
25.0	1
32.0	1

Name: count, dtype: int64

In [9]:


```
df['dependents']=df['dependents'].fillna(0)
df['occupation']=df['occupation'].fillna('self_employed')
df['city']=df['city'].fillna(1080)
df['days_since_last_transaction']=df['days_since_last_transaction'].fillna(0)
```

In [10]: ▶ df

Out[10]:

	customer_id	vintage	age	gender	dependents	occupation	city	customer_nw
0	1	3135	66	1.0	0.0	self_employed	187.0	
1	2	310	35	1.0	0.0	self_employed	1080.0	
2	4	2356	31	1.0	0.0	salaries	146.0	
3	5	478	90	-1.0	0.0	self_employed	1020.0	
4	6	2531	42	1.0	2.0	self_employed	1494.0	
...	...	...	...	...	...	...	...	...
28377	30297	1845	10	0.0	0.0	student	1020.0	
28378	30298	4919	34	0.0	0.0	self_employed	1046.0	
28379	30299	297	47	1.0	0.0	salaries	1096.0	
28380	30300	2585	50	1.0	3.0	self_employed	1219.0	
28381	30301	2349	18	1.0	0.0	student	1232.0	

28382 rows × 21 columns




In [11]: ▶ df = pd.concat([df, pd.get\_dummies(df['occupation'], prefix = str('occupat'), df

Out[11]:

	customer_id	vintage	age	gender	dependents	occupation	city	customer_nw
0	1	3135	66	1.0	0.0	self_employed	187.0	
1	2	310	35	1.0	0.0	self_employed	1080.0	
2	4	2356	31	1.0	0.0	salaries	146.0	
3	5	478	90	-1.0	0.0	self_employed	1020.0	
4	6	2531	42	1.0	2.0	self_employed	1494.0	
...	...	...	...	...	...	...	...	...
28377	30297	1845	10	0.0	0.0	student	1020.0	
28378	30298	4919	34	0.0	0.0	self_employed	1046.0	
28379	30299	297	47	1.0	0.0	salaries	1096.0	
28380	30300	2585	50	1.0	3.0	self_employed	1219.0	
28381	30301	2349	18	1.0	0.0	student	1232.0	

28382 rows × 26 columns



## Scaling Numerical Features for logistic Regression

**since there a lot of outliers in dataset(prev,current bal).distrubutions are skewed so to deal it (Log transformation,Standard Scaler)**

```
In [13]: ▶ num_cols = ['customer_nw_category', 'current_balance',
                        'previous_month_end_balance', 'average_monthly_balance_prevQ2',
                        'current_month_credit', 'previous_month_credit', 'current_month',
                        'previous_month_debit', 'current_month_balance', 'previous_month']

for i in num_cols:
    df[i] = np.log(df[i] + 17000)

std = StandardScaler()
scaled = std.fit_transform(df[num_cols])
scaled = pd.DataFrame(scaled, columns=num_cols)
```

## Model Evaluation'

Recall(to know wrongly marked)

AUC(area under curve),(X axis, ROC (reciever operating curve)

```
In [14]: ▶ df_df_og = df.copy()
df = df.drop(columns = num_cols,axis = 1)
df = df.merge(scaled,left_index=True,right_index=True,how = "left")
```


```
In [15]: ▶ y_all = df.churn
df = df.drop(['churn','customer_id','occupation'],axis = 1)
```

```
In [16]: ▶ baseline_cols = ['current_month_debit', 'previous_month_debit', 'current_b',
                           , 'occupation_retired', 'occupation_salaried', 'occupation_']
df_baseline = df[baseline_cols]
```

```
In [19]: ▶ from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(df_baseline,y_all,test_si
```

```
In [20]: ▶ model = LogisticRegression()
model.fit(xtrain,ytrain)
pred = model.predict_proba(xtest)[: ,1]
```

```
In [21]: ▶ from sklearn.metrics import roc_curve
fpr, tpr, _ = roc_curve(ytest, pred)
auc = roc_auc_score(ytest, pred)
plt.figure(figsize=(12,8))
plt.plot(fpr, tpr, label="Validation AUC-ROC="+str(auc))
x = np.linspace(0, 1, 1000)
plt.plot(x, x, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc=4)
plt.show()
```




## Confusion Matrix

```
In [23]: ▶ pred_val = model.predict(xtest)
label_preds = pred_val

cm = confusion_matrix(ytest,label_preds)

def plot_confusion_matrix(cm, normalized=True, cmap='bone'):
    plt.figure(figsize=[7, 6])
    norm_cm = cm
    if normalized:
        norm_cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        sns.heatmap(norm_cm, annot=cm, fmt='g', xticklabels=['Predicted: No', 'Predicted: Yes'])
    plot_confusion_matrix(cm, ['No', 'Yes'])
```



```
In [24]: ▶ recall_score(ytest,pred_val)
```

Out[24]: 0.13234455219623503

```
In [25]: ▶ def cv_score(ml_model, rstate = 12, thres = 0.5, cols = df.columns):  
    i = 1  
    cv_scores = []  
    df1 = df.copy()  
    df1 = df[cols]  
  
    # 5 Fold cross validation stratified on the basis of target  
    kf = StratifiedKFold(n_splits=5, random_state=rstate, shuffle=True)  
    for df_index, test_index in kf.split(df1, y_all):  
        print('\n{} of kfold {}'.format(i, kf.n_splits))  
        xtr, xvl = df1.loc[df_index], df1.loc[test_index]  
        ytr, yvl = y_all.loc[df_index], y_all.loc[test_index]  
    baseline_scores = cv_score(LogisticRegression(), cols = baseline_cols)
```

1 of kfold 5

1 of kfold 5

1 of kfold 5

1 of kfold 5

1 of kfold 5

In [ ]: ▶